

Shikhar Shiromani

AI Safety and Interpretability Researcher

✉ rbk.shikhar@gmail.com 📍 San Diego, CA 🗣 [shikhar1729](#) 🌐 [shikhar-shiromani](#) 📄 [Google Scholar](#)

Summary

AI safety and interpretability researcher working on chain-of-thought faithfulness, steganographic reasoning, and the monitorability of large language models. Machine Learning Engineer at Apple; Pivotal Research Fellow and MARS V Research Fellow.

Research Fellowships

Pivotal Research Fellow

Feb 2026 – present

Chain-of-thought faithfulness in LLMs, Pivotal Research (London)

- Selected as a Research Fellow from 2,500+ applicants; awarded a merit-based extension in March 2026.
- Working with Noah Y. Siegel (Google DeepMind, UK) on chain-of-thought faithfulness.
- Showed CoT monitors collapse from ~95% to under 11% catch rate on the subset where the reasoning trace is the only signal, and that the attack transfers across monitor families and agent models.

MARS V Research Fellow

July 2026 – present

Steganographic reasoning in LLMs, Meridian (Cambridge, UK)

- Selected as one of 86 fellows from 1,733 applications (approximately 5%).
- Working with Julian Schulz on how LLMs hide chain-of-thought in text.
- Evidence for a distributed mechanism that transfers across encoding schemes by reusing the same weights.

SPAR Fellow

Feb 2026 – May 2026

Attention Consistency Training for Safety, Kairos

- Accepted from a round of nearly 2,500 applications.

Research

A False Average - CoT Monitor Collapse

May 2026 – Aug 2026

arXiv 2608.00583

- Gradient-free rewrite of an agent's reasoning, holding every command and output fixed, drops a held-out monitor's catch rate from ~95% to under 11% on CoT-only hacks.
- Demonstrated that aggregate monitor accuracy is a false average that hides near-total collapse on the subset monitors exist to cover.
- Trace-only defenses recover it only partially; a probe on an open-weight surrogate monitor's activations separates missed hacks, with a causal control showing this is a detector rather than latent monitor knowledge.

The Hypocrisy Gap - Internal Belief vs. CoT via Sparse Autoencoders

Nov 2025 – Jan 2026

arXiv 2602.02496

- Introduced a mechanistic metric quantifying divergence between a model's internal belief and its generated explanation using sparse autoencoders and sparse linear probes.
- AUROC 0.55-0.74 for detecting sycophantic and hypocritical runs on Gemma, Llama, and Qwen, consistently outperforming a decision-aligned log-probability baseline (0.41-0.50).

Benchmarking Pluralistic Alignment (PAB) ICML 2026 Pluralistic Alignment Workshop • Built the Pluralistic Alignment Benchmark covering 320 scenarios, 8 behavioral-risk categories, and 63 synthetic personas. • Evaluated 9 leading LLMs with a multi-judge pipeline, revealing substantial persona-conditioned drift in manipulation, concealment, persuasion, and sycophancy.	Mar 2026 – June 2026
VLM Reliability Probe ICLR 2026 Multimodal Intelligence and Multimodal Reasoning Workshops • Cross-family reliability probe testing whether spatial attention structure predicts VLM correctness. • Attention coherence has near-zero correlation with accuracy ($R \approx 0.001$); self-consistency is far stronger ($R = 0.43$). • Causal ablations expose a split between late-fusion and early-fusion architectures with direct monitor-design implications.	Dec 2025 – June 2026
COMPASS - Context-Modulated PID Attention Steering AAAI 2026 LaMAS Workshop (Oral), arXiv 2511.14776 • Interpretable feedback-driven decoding framework steering attention via a Context Reliance Score. • Reduced contextual hallucinations by 2.8-5.8 points across HotpotQA, XSum, HaluEval, and RAGTruth without retraining.	Aug 2025 – Nov 2025
Linear Predictability of Attention Heads arXiv 2603.13314 • Revealed emergent inter-head linear structure enabling cross-head QKV approximation from 2-5 reference heads. • Showed the structure is learned rather than architectural, emerging during pretraining, supported by a theoretical lower bound at initialization. • Compressed the KV cache via learned projections for 2x reduction with minimal performance loss.	Aug 2024 – May 2025
ChameleonBench - Quantifying Alignment Faking ACML 2025 (Oral) • 800-prompt benchmark across 8 misbehavior categories with an external judge pipeline. • Across 6 leading LLMs, up to 20% higher harmful response rates in free-form versus test-like settings.	May 2025 – June 2025
ProMoral-Bench - Prompting Strategies for Moral Reasoning and Safety NeurIPS 2025 PersonaNLP, arXiv 2602.13274 • Benchmark comparing 11 prompting strategies across 4 LLM families, including the new ETHICS-Contrast dataset. • Introduced the Unified Moral Safety Score; exemplars outperform multi-hop reasoning on accuracy, safety, and efficiency.	May 2025 – Nov 2025
SAGE - Agentic Framework for Pathology Biomarker Discovery arXiv 2602.00953 • Agentic framework discovering pathology biomarkers grounded in biological literature and multimodal evidence. • Links image-derived features with molecular markers and clinical outcomes to improve clinical translatability.	July 2025 – present

Experience

Apple, Machine Learning Engineer

- Developing intelligent systems that optimize power and performance across Apple's operating systems.

San Diego, CA
June 2026 – present
4 months

NVIDIA, Senior Software Engineer, Tech Lead

- Built the first server lifecycle management system at NVIDIA, streamlining resource requests, allocations, and deployment tracking for faster data center rack bring-up.
- Explored ML techniques to enhance node uptime by predicting failures and optimizing system reliability.

Santa Clara, CA
Feb 2025 – June 2026
1 year 5 months

NVIDIA, Machine Learning Research Intern

- Architected an ML-driven predictive maintenance system for 1000+ Blackwell and Hopper nodes, delivering over \$100k/year in operational savings.
- Implemented multidimensional multimodal anomaly detection with negative sampling and integrated gradients for variable attribution during model interpretation.
- Built microservices for telemetry collection (6200+ metrics in 2 minutes), Grafana dashboarding, and failure prediction.
- Deployed an Airflow MLOps pipeline and a FastAPI REST service handling 10,000+ inferences/day.
- Presented a poster at NTech Conference 2024. Filed for patent as first author.

Santa Clara, CA
May 2024 – Aug 2024
4 months

AI Institute for Advances in Optimization (AI4OPT), Georgia Tech, Graduate Research Assistant, advised by Prof. Pascal Van Hentenryck

- Founding backend engineer for a highly available on-demand multi-modal transportation system focused on large-scale ride-sharing, shipped on iOS and Android.
- Piloted in Savannah, GA over 15 months, scheduling 100+ rides daily.
- Project awarded one of the SMART 20 Awards 2025.

Atlanta, GA
Aug 2023 – Dec 2024
1 year 5 months

NVIDIA, Software Engineer II

- Architected a scalable self-serve CI/CD platform, cutting engineering workload turn-around time by 93%.
- Presented a paper at NTech India Conference on the benefits of the system.
- Delivered the NvCCP (NVIDIA Camera Control Protocol) module for automotive safety assessments, driving software architecture, features, and functional verification.

Bangalore, India
July 2019 – July 2023
4 years 1 month

Education

Georgia Institute of Technology, Computational Science and Engineering

Atlanta, GA, USA
Aug 2023 – Dec 2024

Birla Institute of Technology and Science, Pilani, Electrical and Electronics Engineering

Pilani, India
Aug 2015 – June 2019

Service and Mentorship

Peer Reviewer

Sept 2025 – present

12 completed reviews recorded on OpenReview

- Efficient Reasoning Workshop, NeurIPS 2025 (3 reviews) and COLM 2026 (5 reviews, repeat invitation from the same organizing committee).
- BlackboxNLP 2026 at EMNLP 2026, Budapest (4 reviews), nominated by the program chair committee.
- Reviewer pool, InterpScience 2026 at NeurIPS 2026; reviewer invitation, Pluralistic Alignment Workshop at ICML 2026.

Research Mentor, Algoverse

July 2025 – present

ML Researcher and Summer 2026 mentor

- Invited back as a research mentor for the Summer 2026 cohort.

IEEE Senior Member

June 2026 – June 2026

Elevated June 2026

- Highest IEEE grade available by application, held by approximately 10% of members.

Patent Pending (first author)

Mar 2025 – Mar 2025

Machine Learning-Based Anomaly Detection Using Data Center Metrics

- Filed March 2025, developed at NVIDIA against Blackwell and Hopper node telemetry.

Technical Skills

Languages: Python, C/C++, TypeScript, JavaScript, SQL**ML and Interpretability:** PyTorch, sparse autoencoders, linear probes, activation patching, integrated gradients, TransformerLens-style analysis**Systems and Infrastructure:** Kubernetes, Docker, Airflow, FastAPI, Django/Flask, NodeJS, ReactJS, MySQL, InfluxDB, Jenkins

Awards

Pivotal Research Fellowship: Selected from 2,500+ applicants, plus a merit-based extension (2026)**MARS V Fellowship:** 86 selected from 1,733 applications, approximately 5% (2026)**SPAR Fellowship:** Kairos, from a round of nearly 2,500 applications (2026)**IEEE Senior Member:** Elevated June 2026**SMART 20 Awards 2025:** AI4OPT on-demand multi-modal transportation project**Competitively selected orals:** ACML 2025 (ChameleonBench), LaMAS at AAAI 2026 (COMPASS), SELVA at ACL 2026 (Linear Predictability, one of four orals)